EXPANDED UK & IRELAND MSP PERFORMANCE BENCHMARK

The MSP Metrics Playbook

A deeper, more practical guide to the financial, operational, customer, people, security and strategic metrics that separate high-performing managed service providers from the rest.

Expanded vendor-neutral edition
Designed to accompany the MSP Benchmark Scorecard.

May 2026
GBP-standardised where public benchmark data was originally USD.

EXECUTIVE SUMMARY

What top-performing MSPs measure differently

The strongest MSPs do not win because they measure everything. They win because they consistently measure the few numbers that expose pricing quality, delivery leverage, recurring revenue quality, client retention and security maturity. The pattern is consistent: top performers are usually better priced, more standardised, more automated, less project-dependent and more deliberate about account expansion.

19%+

Adjusted EBITDA is a common best-in-class threshold in MSP benchmark material.

42%+

Blended gross margin signals whether pricing and delivery can support strong profit.

75%+

Recurring revenue is a practical threshold for M&A-grade; predictability.

£100-£150+

Premium UK/Ireland PUPM suggests a deeper managed service relationship.

200+

Endpoints per technician indicates mature standardisation and automation.

<5%

Annual logo churn reflects a sticky, well-managed client base.

How this expanded edition is structured

Part 1 expands the 12 metrics used by the public scorecard. Each metric includes a definition, benchmark bands, how to interpret it, useful segmentation and practical actions.

Read this as a management tool

The goal is not a perfect academic benchmark. The goal is to give MSP owners and leadership teams a clear operating language for better pricing, cleaner delivery, stronger retention and more profitable growth.

Part	What it covers
Part 1	The 12 scorecard metrics in detail.
Part 2	Additional operating metrics for mature MSPs.
Part 3	UK and Ireland market context, regulation and valuation.
Part 4	Practical playbooks, 90-day plan and methodology notes.

THE OPERATING MODEL

The logic behind the metrics

Most MSP performance issues are connected. Weak pricing depresses gross margin. Weak gross margin limits EBITDA. Low PUPM usually means services are missing or being given away. Missing services reduce wallet share and can increase risk. Poor standardisation increases tickets, tickets consume technicians, and overloaded technicians damage client experience.

Financial layer

Gross margin, EBITDA, recurring percentage, PUPM and average client show whether the comme healthy.

Delivery layer

Revenue per employee, revenu technician, endpoints per techn tickets per endpoint show whet can scale service delivery.

Client layer

Logo churn, average services per client and security attach rate show whether the relationship is sticky, deep and risk-aware.

The main idea

A benchmark score is useful only if it leads to a decision: reprice, standardise, automate, consolidate, review, expand, exit or invest.

Layer	Common symptom	Likely decision
Pricing	PUPM below target	Repackage, reprice, remove free scope
Delivery	Tickets or reactive hours too high	Standardise, automate, fix recurring issues
Retention	Churn or low QBR attendance	Improve value reviews and roadmap discipline
Security	Low attach rate or baseline gaps	Define non-negotiable controls and adoption plan

01

PROFITABILITY DISCIPLINE

Blended gross margin

Formula: (Revenue - COGS) / Revenue

Definition: The percentage of revenue left after direct delivery costs. For an MSP, COGS should include front-line technical labour, service or project management where directly tied to delivery, tooling, third-party licences, subcontractors and direct cloud or service costs.

Top-quartile	Healthy / median	Watch zone
>42% blended	**35-38%**	**Below 30%**

Why this matters

Gross margin is the clearest early signal of whether the delivery model can produce strong EBITDA. If gross margin is weak, overhead cuts rarely fix the business. The root cause is usually pricing, labour efficiency, tool recovery, client standards or too much bespoke service effort.

How to read it

- High gross margin with weak growth may indicate an efficient but under-invested business.
- Low gross margin with strong growth can look impressive in revenue terms but may be scaling operational pain.
- Service labour must be included. Excluding it makes the metric look better while hiding delivery leakage.

Useful segmentation	Benchmark / note
Managed services	50-60%+ target for mature operators
Projects	35-50% depending on scope control
Product resale	15-20% is common; avoid diluting the model
Managed security	Often 55%+ when packaged and standardised

Practical improvement moves

- Split margin by managed services, projects, resale, cloud and security.
- Review contracts where effective hourly rate is below target.
- Recover tool costs explicitly through packaged pricing.
- Standardise onboarding and remove unsupported exceptions.
- Reprice legacy agreements where scope expanded without uplift.

02

VALUATION ENGINE

EBITDA / net profit margin

Formula: Adjusted EBITDA / Revenue

Definition: Operating profit before interest, tax, depreciation and amortisation, adjusted for owner compensation and one-off items. Smaller founder-led MSPs should normalise owner pay to make the number comparable.

Top-quartile	Healthy / median	Watch zone
19%+ adjusted EBITDA	**11-15%**	**Below 8-10%**

Why this matters

EBITDA is the metric lenders, acquirers and serious operators focus on. It shows whether the MSP converts revenue into cash-generating operating profit. The highest quality MSPs combine margin with recurring revenue, low churn and low customer concentration.

How to read it

• EBITDA below 10% does not always mean the business is weak, but it demands a clear explanation.

• A growing MSP may temporarily accept lower EBITDA if spend is deliberate and growth-efficient.

• Owner compensation, family payroll and one-off expenses should be normalised before comparison.

Useful segmentation	Benchmark / note
Small MSP	5-12% is common, often distorted by owner pay
Mid-market	12-18% indicates a healthier platform
Large / PE-grade	15-22%+ with recurring revenue and process maturity
M&A-grade;	20%+ is a strong signal when quality is clean

Practical improvement moves

• Fix gross margin before cutting useful sales or client-success investment.

• Segment clients by contribution margin and ticket noise.

• Separate true growth investment from poor operating discipline.

• Build a monthly management pack showing margin, churn and recurring revenue.

• Use a rolling 12-month view to avoid reacting to one project-heavy month.

03

REVENUE QUALITY

Recurring revenue percentage

Formula: Annualised recurring revenue / Total annual revenue

Definition: The proportion of total annual revenue that comes from recurring contracts, subscriptions or managed service agreements rather than one-off projects, hardware or ad hoc labour.

Top-quartile	Healthy / median	Watch zone
75%+	**50-70%**	**Below 50%**

Why this matters

Recurring revenue creates predictability, improves planning, reduces reliance on project timing and usually receives a stronger valuation multiple. Project revenue can be valuable, but a project-heavy MSP is harder to forecast and often more labour constrained.

How to read it

- A high recurring percentage is not enough if agreements are underpriced or churn-prone.
- Project revenue is healthier when it feeds recurring services, not when it replaces them.
- Recurring revenue should include clear scope, not unlimited support disguised as MRR.

Useful segmentation	Benchmark / note
Recurring dominant	75-85%+
Mixed model	50-70%
Project heavy	Below 50%
Investor lens	Quality can matter as much as absolute revenue

Practical improvement moves

- Create standard managed service packages and move clients away from ad hoc support.
- Attach security, backup and compliance services to the recurring base.
- Set minimum monthly recurring commitments for new clients.
- Track revenue by recurring, project, resale and cloud.
- Avoid using hardware refreshes to mask weak recurring growth.

04

PRICING POWER

AISP / PUPM

Formula: Monthly recurring revenue / Supported users

Definition: All-in seat price, also called per-user-per-month. This shows how much recurring revenue the MSP earns for each supported user after bundling support, security, backup, strategic service and other recurring components.

Top-quartile	Healthy / median	Watch zone
£100-£150+ UK/Ireland	**£65-£95**	**Below £65**

Why this matters

PUPM is one of the most actionable pricing metrics. It quickly shows whether the MSP is selling a proper managed service stack or a thin support layer. Two MSPs can both support 500 users, but the one at £120 PUPM has far more room to deliver security, proactive work and profit than the one at £55.

How to read it

• Low PUPM usually means security, backup, documentation or strategic work is being sold separately or not sold at all.

• High PUPM must be supported by clear value, proactive service and visible outcomes.

• Regulated verticals such as finance, healthcare and legal can justify higher PUPM when service scope is real.

Useful segmentation	Benchmark / note
Basic support	£40-£65
Standard SME	£65-£95
Premium / regulated	£100-£150+
Enterprise / co-managed	£150+ depending on responsibility

Practical improvement moves

• Create a client-by-client PUPM ranking.

• Build a clear minimum service baseline.

• Use good-better-best packaging to avoid bespoke quoting.

• Show clients which risks are included and which are outside scope.

• Reprice slowly but deliberately, starting with renewals and low-margin accounts.

05

CLIENT QUALITY

Average MRR per client

Formula: Total monthly recurring revenue / Recurring clients

Definition: The average monthly recurring revenue generated by each client. It is a simple way to understand whether the client base is large enough to support proactive service, account management and profitable delivery.

Top-quartile	Healthy / median	Watch zone
£8k-£20k+ for larger MSPs	**£1.5k-£4k SMB range**	**Below £1k-£1.5k**

Why this matters

Very small clients often consume disproportionate effort. They may create tickets, admin time, billing complexity and security risk without contributing enough MRR to fund proactive work. A higher average MRR per client usually gives the MSP room to deliver properly.

How to read it

• A low average can be fine if the MSP has an extremely standardised, low-touch model.

• The danger is a low-MRR client that still expects bespoke, high-touch service.

• Average MRR should be analysed alongside ticket volume, margin and churn.

Useful segmentation	Benchmark / note
Small MSP	£800-£2k per client can be common
Mid-market MSP	£3k-£8k
Larger accounts	£10k-£50k+
Strategic concern	Low MRR plus high noise is the worst combination

Practical improvement moves

• Set a minimum MRR threshold for new clients.

• Create a plan for legacy micro-accounts: uplift, reduce scope or exit.

• Bundle more of the client roadmap into recurring agreements.

• Track account management time spent per client, not just support tickets.

• Avoid taking on small clients that do not fit technical standards.

06

WHOLE-BUSINESS PRODUCTIVITY

Revenue per employee

Formula: Annual revenue / Total FTEs

Definition: A broad productivity measure across the whole business, including service, sales, leadership, finance, admin and operations. It shows whether the company scales with process or simply adds headcount whenever revenue grows.

Top-quartile	Healthy / median	Watch zone
£120k+ per employee	**£80k-£110k**	**Below £70k**

Why this matters

Revenue per employee captures the combined effect of pricing, service design, automation, sales productivity and overhead structure. It is especially useful for comparing similar-sized MSPs, but should not be used in isolation because revenue mix matters.

How to read it

• A project-heavy MSP can show high revenue per employee while hiding weak recurring quality.

• A growing MSP may temporarily show lower RPE while hiring ahead of demand.

• Compare the trend over time more than a single month.

Useful segmentation	Benchmark / note
Lower leverage	Below £70k
Typical	£80k-£110k
Strong	£120k+
Very strong	£160k+ where pricing and automation are mature

Practical improvement moves

• Analyse revenue per department, not just total company RPE.

• Reduce manual admin with clean PSA, billing and documentation workflows.

• Make sure non-technical roles improve sales, retention or delivery leverage.

• Remove services that require too much bespoke labour.

• Use automation to keep revenue growing faster than headcount.

07

SERVICE LEVERAGE

Revenue per technician

Formula: Annual revenue / Technical delivery FTEs

Definition: The amount of annual revenue supported by each service-delivery employee. This focuses specifically on the part of the organisation that is usually the constraint in an MSP: technical capacity.

Top-quartile	Healthy / median	Watch zone
£170k+ per service FTE	**£140k-£170k**	**Below £140k**

Why this matters

Technicians are expensive, hard to hire and central to service quality. If each technician supports too little revenue, the MSP either needs better pricing, less reactive noise, more automation or clearer service scope.

How to read it

• Very high revenue per technician can indicate leverage, but it can also indicate understaffing if SLA performance and burnout are poor.

• Low revenue per technician usually points to underpriced agreements, too many tickets or excessive bespoke work.

• Look at revenue per technician beside utilisation and ticket reopen rates.

Useful segmentation	Benchmark / note
Support-heavy	Below £140k
Healthy	£140k-£170k
Strong	£170k+
Check quality	High leverage only counts if service outcomes remain good

Practical improvement moves

• Separate L1, L2, L3 and project capacity where possible.

• Track reactive hours and project hours separately.

• Improve knowledge base and escalation paths.

• Reduce low-value work with self-service and automation.

• Use client standards to reduce preventable issues.

08

AUTOMATION MATURITY

Endpoints per technician

Formula: Managed endpoints / Technical delivery FTEs

Definition: The number of endpoints supported per technician. It is a practical measure of standardisation, automation maturity, client environment quality and support model efficiency.

Top-quartile	Healthy / median	Watch zone
200-300+	**100-180**	**Below 75-100**

Why this matters

MSPs with mature RMM automation, clean documentation and consistent client standards can support more endpoints per technician without service quality collapsing. Low ratios often reveal noisy environments, weak onboarding or tools that are not being fully used.

How to read it

- Do not compare endpoint-heavy and user-heavy models blindly.
- Co-managed contracts can lift endpoint ratios because internal IT handles some frontline demand.
- Endpoint count should be paired with tickets per endpoint and reactive hours per endpoint.

Useful segmentation	Benchmark / note
Warning	Below 75
Average	100-ish
Upper quartile	150-200
World class	200-300+ when quality remains strong

Practical improvement moves

- Create standard device, patching and security baselines.
- Use RMM scripts for the top recurring fixes.
- Improve onboarding so new clients are normalised early.
- Retire unsupported legacy systems or price them as exceptions.
- Measure reactive hours per endpoint monthly.

09

OPERATIONAL NOISE

Tickets per endpoint per month

Formula: Monthly tickets / Managed endpoints

Definition: The number of support tickets generated per endpoint each month. This is one of the clearest indicators of preventable operational noise inside the client base.

Top-quartile	Healthy / median	Watch zone
Below 0.5, best around 0.25	**0.5-1.0**	**Above 1.0**

Why this matters

Ticket volume consumes technical capacity, damages customer experience and makes scaling difficult. A business can grow MRR while quietly drowning the service desk if ticket density is rising.

How to read it

- High ticket volume is not always bad if it reflects a short-term onboarding project or a large transition.
- Sustained high volume usually means weak standards, poor user training, recurring issues or missing automation.
- Tickets should be categorised by avoidable, user-generated, incident, request and project-related demand.

Useful segmentation	Benchmark / note
World class	~0.25 tickets per endpoint per month
Good	Below 0.5
Average	0.5-1.0
Watch	Above 1.0

Practical improvement moves

- Build a recurring issues dashboard.
- Review top clients by ticket density, not just total tickets.
- Create a top-10 automation backlog.
- Use knowledge base articles and user education for repetitive requests.
- Link noisy accounts to pricing and service-review conversations.

10

RETENTION HEALTH

Annual logo churn

Formula: Clients lost during period / Clients at start of period

Definition: The percentage of clients that leave over a 12-month period, regardless of revenue value. This measures relationship retention at the account level.

Top-quartile	Healthy / median	Watch zone
Below 5%	**8-12%**	**Above 15-20%**

Why this matters

Logo churn damages growth, creates replacement pressure and undermines valuation quality. MSP churn can be deceptively quiet because unhappy clients may stay until renewal, procurement review or a major incident forces a change.

How to read it

• Low churn with low PUPM may mean clients are sticky because the MSP is undercharging.

• High-margin clients can churn because they experience few visible issues and undervalue prevention.

• Revenue churn and logo churn should both be tracked.

Useful segmentation	Benchmark / note
Top quartile	Below 5%
Typical mature MSP	8-12%
Warning	15-20%+
Expansion lens	NRR can be strong even with some logo churn if retained clients expand

Practical improvement moves

• Run value reviews with decision-makers, not only technical contacts.

• Track review attendance, roadmap acceptance and unresolved complaints.

• Document avoided incidents and completed improvements.

• Analyse churn by client size, vertical, service tier and salesperson.

• Address low-MRR, high-noise clients proactively.

WALLET SHARE

11 Average services per client

Formula: Total recurring services sold / Active clients

Definition: The average number of recurring services each client buys. Services may include helpdesk, RMM, security, backup, BCDR, email security, identity protection, compliance support, documentation, strategic reviews and cloud management.

Top-quartile	Healthy / median	Watch zone
6+ services	**3-4**	**1-2**

Why this matters

Service depth indicates how much of the client's IT risk and operating environment the MSP owns. More services usually mean higher wallet share, stronger retention and better economics, provided the services are standardised and profitable.

How to read it

• More services is not automatically better if each one is bespoke and poorly adopted.

• A low service count may indicate white space, but it may also reflect a deliberately narrow specialist model.

• Measure service adoption alongside gross margin and client satisfaction.

Useful segmentation	Benchmark / note
Thin relationship	1-2 services
Average	3-4 services
Deep relationship	6+ services
Strategic goal	Move from tool sale to business outcome ownership

Practical improvement moves

• Build a white-space matrix across the client base.

• Define which services belong in the minimum standard.

• Use QBRs to connect service gaps to business risks.

• Prioritise security, backup and resilience before nice-to-have add-ons.

• Track attach rates by account manager and client segment.

12

RISK COVERAGE

Security attach rate

Formula: Clients using managed security services / Total clients

Definition: The percentage of clients using security services such as EDR/MDR, email security, phishing training, vulnerability management, backup, BCDR, identity monitoring or compliance support. The exact definition should be consistent internally.

Top-quartile	Healthy / median	Watch zone
85-100%	**50-70%**	**Below 50%**

Why this matters

Security attach rate is a commercial metric and a risk metric. If clients are poorly protected, the MSP may still be blamed when incidents happen. A strong security baseline reduces risk, improves recurring revenue and strengthens the MSP's advisory position.

How to read it

- Define what counts as security attach. Basic AV alone should not be treated the same as a complete managed security baseline.
- Attach rate should be measured by service type, not only at an aggregate level.
- Regulated clients should have a higher baseline than low-risk clients.

Useful segmentation	Benchmark / note
Weak	Below 50%
Developing	50-70%
Strong	85%+
Best practice	Security baseline is part of the managed service, not an optional afterthought

Practical improvement moves

- Create a minimum security standard for every client.
- Run a gap analysis by account.
- Use risk language rather than product language.
- Track adoption of backup, MFA, endpoint protection, email security and vulnerability separately.
- Escalate clients who refuse baseline controls and document the risk.

PART 2

The next layer of MSP metrics

The 12 scorecard metrics tell you where the MSP stands. The next metrics explain why. Use them once the core data is stable enough to support management decisions.

Recommended sequence

Start with gross margin, EBITDA, recurring revenue, PUPM, endpoints per technician, tickets per endpoint and churn. Then add effective rate, service multiple of wages, RHEM, NRR, DSO and customer concentration.

Pricing and margin diagnostics

Metric	Interpretation and use
Effective rate Monthly recurring contract value divided by the hours delivered against that contract.	Reveals whether a fixed-fee agreement is economically healthy. A client paying £3,000 per month but consuming 30 hours has an effective rate of £100/hour. That may be too low once labour, tools, overhead and risk are considered.
Contribution rate Client revenue less labour and direct product/tool cost, divided by hours worked.	Useful when client agreements contain a mix of labour, licences and project effort. It helps identify which clients contribute profit per technical hour and which clients consume capacity without enough return.
Service multiple of wages Service revenue divided by service wages.	A compact productivity index. Mature operators often target 2.5x+ and can approach 3.0x. Below 1.9x suggests delivery economics need attention.
Gross margin by service line Margin split across managed services, projects, resale, cloud and security.	Blended GM can hide weak service lines. Project revenue, resale or unmanaged cloud can dilute the model if priced poorly.

Practical use

These metrics are best reviewed at client or agreement level. They help explain why the P&L; looks the way it does. If blended margin is weak, these diagnostics usually reveal the contracts causing the issue.

Retention and expansion metrics

Metric	Interpretation and use
Net revenue retention Starting client revenue plus expansion, less contraction and churn, divided by starting revenue.	One of the most important account-management metrics. Above 110% means the existing base expands even before new-logo sales.
Gross revenue retention Starting client revenue less contraction and churn, before expansion.	Shows the durability of existing revenue before upsell masks the issue. A healthy base usually holds 90%+ before expansion.
Revenue churn Recurring revenue lost during the period divided by opening recurring revenue.	Useful where losing one large client matters more than losing several small clients. Track separately from logo churn.
QBR coverage Percentage of strategic clients receiving timely review meetings with proper decision-maker attendance.	Turns operational delivery into visible business value. Strong QBRs create expansion, reduce silent churn and align budgets to roadmaps.

Practical use

Retention should be viewed through both logo and revenue lenses. A small client leaving is not the same as a strategic client leaving, and expansion can sometimes hide shrinking trust in parts of the base.

ADVANCED METRICS

Service delivery metrics

Metric	Interpretation and use
Reactive hours per endpoint per month Reactive support hours divided by endpoints.	A leading indicator for margin. World-class operators push toward 15 minutes or less. Noisy estates can sit at 45-60 minutes.
First response time Time from ticket creation to first meaningful response.	Important for customer experience, but should not be optimised at the expense of resolution quality. Segment by priority.
Mean time to resolution Average time to close tickets or incidents.	Useful when segmented by severity. A single average can be misleading if password resets and major incidents are blended together.
Reopen rate Percentage of closed tickets reopened by users or technicians.	A practical quality indicator. High reopen rates suggest rushed fixes, poor root-cause analysis or weak communication.

Practical use

Service metrics should be segmented by priority, client, service type and root cause. A single average can create false comfort.

ADVANCED METRICS

Cash, sales and resilience metrics

Metric	Interpretation and use
DSO Accounts receivable divided by credit revenue, multiplied by 365.	Long collection cycles starve growth cash. Autopay or direct debit, clean invoicing and prompt follow-up can materially improve cash conversion.
CAC payback Sales and marketing cost to acquire a client divided by gross profit from that client.	Shows whether growth is efficient. MSP-specific data is patchy, but below 12 months is strong and above 24 months deserves scrutiny.
Customer concentration Revenue share from largest client, top five and top ten clients.	A top client above 15% or top five above 40% can create valuation and resilience risk.
Rule of 40 Revenue growth percentage plus EBITDA margin percentage.	A high bar for recurring-revenue businesses. Useful as a strategic lens, but not every good MSP needs to hit it every year.

Practical use

Cash and resilience metrics are often ignored until there is pressure. They should be part of monthly management review because they influence growth capacity and strategic optionality.

Security and compliance metrics

Metric	Interpretation and use
Patch compliance Percentage of endpoints or systems patched within policy timing by severity.	Critical for security and compliance. UK Cyber Essentials requirements make patching discipline commercially visible.
Backup success rate Successful backups divided by scheduled backups.	A high backup success rate is necessary but not sufficient. Restore testing matters more than backup activity alone.
Restore test success Successful restore tests divided by attempted restore tests.	Shows whether resilience is real. If restores are not tested, backup confidence is largely theoretical.
MFA and EDR coverage Percentage of users, admins and endpoints covered by baseline security controls.	Security baselines should be measurable. Coverage gaps create client risk and MSP reputational exposure.

Practical use

Security metrics should be client-visible where appropriate. They are not just internal technical KPIs. They help clients understand risk, evidence and accountability.

PART 3

UK and Ireland commercial context

Pricing bands

Basic managed IT packages commonly sit around £40-£65 PUPM. Standard SME packages often sit around £65-£95. Premium, security-led, co-managed or regulated service models can move into £100-£150+ and sometimes beyond.

Buyer expectations

UK and Irish SMB buyers increasingly expect cyber, backup, compliance evidence and strategic guidance to sit alongside support. The conversation is moving from tickets and response times toward risk, continuity and accountability.

Labour market pressure

Technical talent remains expensive and difficult to hire. This makes automation, documentation and client standardisation more important. Hiring more technicians is rarely the only answer.

Seasonality

Sales cycles can be slower around August, Christmas and procurement-heavy periods. This matters when comparing pipeline conversion or sales-cycle metrics across quarters.

Where original public benchmarks were in USD, GBP ranges are rounded using an approximate mid-2025 planning rate of £1 = $1.27. Treat converted figures as directional, not accounting-grade conversions.

Regulation as a benchmark driver

Regulation is turning security and resilience metrics into commercial requirements. MSPs serving regulated or supply-chain-sensitive clients should expect more scrutiny around patching, identity, backup, restore testing, incident response and evidence.

Area	Why it matters for MSP metrics
Cyber Essentials	Raises the importance of patch compliance, MFA coverage, secure configuration and evidence. It also creates a clear commercial entry point for compliance support.
NIS2	Brings many ICT service providers and essential or important entities into a stronger incident reporting and risk-management regime. MSPs need clearer security baselines and incident processes.
DORA	Relevant to financial entities and ICT third-party risk. It increases scrutiny around resilience, recovery time, contracts, concentration risk and testing.
GDPR / UK data protection	Makes data handling, backup, breach response, access control and documentation commercially visible in client conversations.
Cyber insurance	Insurers increasingly ask about MFA, EDR, backup, awareness training and incident response. Security attach rate can become a client insurability issue.

Valuation lens

The same metrics that make an MSP easier to run also make it easier to finance or sell. Buyers want predictable recurring revenue, strong gross margin, clean contracts, low churn, low client concentration and a management team that can operate beyond the founder.

Positive valuation signals

75%+ recurring revenue. 20%+ EBITDA with clean adjustments. Strong gross margin and service-line profitability. Low churn and strong NRR. Low customer concentration. Standardised contracts and packaged services. Documented processes and management reporting.

Valuation drag factors

Founder dependency. Top client concentration above 15%. Project-heavy revenue mix. Poor tool utilisation. Weak EBITDA after normalisation. Unclear COGS allocation. High churn or unmanaged low-MRR accounts.

Owner takeaway

Even if an exit is not planned, valuation-quality metrics are still useful. They force operational clarity and reveal whether growth is creating enterprise value or just adding complexity.

PART 4

Diagnostic playbook: symptoms, causes and responses

Symptom	Likely root cause	Practical response
Gross margin below 30%	Underpriced agreements, weak COGS allocation, tool cost leakage or excessive manual delivery.	Reprice legacy accounts, include service labour in COGS, standardise scope and recover tool costs through packaged pricing.
EBITDA below 10%	Weak gross margin, bloated overhead, unmanaged client noise or poor revenue mix.	Fix margin first, then review overhead and investment quality. Avoid cutting growth activities that are actually working.
PUPM below £65	Thin support model or security/compliance value not monetised.	Build tiered packages, define minimum security and resilience baselines, then migrate clients over time.
Average MRR per client below £1k	Too many micro-clients or no minimum client profile.	Introduce a floor for new business and create an uplift or exit plan for legacy accounts.
Endpoints per tech below 75	Noisy client environments, weak automation, poor documentation or too many exceptions.	Normalise client standards, automate recurring fixes and measure reactive hours per endpoint.
Tickets per endpoint above 1.0	Preventable issues, poor onboarding, weak patching, user training gaps or unreliable infrastructure.	Create a top recurring-issues list and tie noisy clients to technical standards and pricing conversations.
Logo churn above 15%	Weak perceived value, poor-fit clients, service inconsistency or unmanaged expectations.	Run decision-maker reviews, build client roadmaps, document improvements and track early churn signals.
Security attach below 50%	Security positioned as optional rather than baseline risk management.	Define a minimum control set and show gaps in plain business language.
DSO above 45 days	Loose billing process, no autopay/direct debit or slow collections follow-up.	Mandate autopay for recurring services, clean invoice data and follow missed payments quickly.
Top client above 15% of revenue	Concentration risk or over-reliance on a founder-era account.	Build a diversification plan and avoid customising the whole operating model around one client.

90-day implementation plan

Days 1-15: baseline

Calculate the 12 scorecard metrics using best-available data. Document how each metric was calculated. Agree which numbers are reliable and which are estimates. Separate recurring, project, resale and cloud revenue.

Days 16-30: segment

Rank clients by PUPM, MRR, ticket density and margin. Identify noisy, underpriced and strategically important accounts. Build a white-space matrix for services per client. Spot immediate security baseline gaps.

Days 31-60: decide

Select three improvement priorities. Create a repricing or packaging plan. Build a top recurring-issues automation backlog. Prepare review agendas for key clients.

Days 61-90: act

Run targeted client reviews. Move selected clients to new packages or standards. Implement the highest-impact automations. Update the scorecard and compare trend movement.

Keep the cadence simple

Monthly is enough for most metrics. Weekly can help during service desk improvement projects, but too many dashboards can create noise. The best metric cadence is one that creates decisions.

Suggested management dashboard

Dashboard section	Metrics	Management question
Profitability	Gross margin, EBITDA, service-line margin, contribution rate.	Are we converting revenue into profit?
Revenue quality	Recurring revenue %, average MRR per client, PUPM, NRR, GRR.	Is the base predictable and expanding?
Delivery leverage	Endpoints per technician, revenue per technician, tickets per endpoint, RHEM, utilisation.	Can service scale without simply adding people?
Client health	Logo churn, revenue churn, NPS/CSAT, QBR coverage, services per client.	Are clients seeing enough value to stay and expand?
Security baseline	Security attach, MFA coverage, EDR coverage, patch compliance, backup and restore testing.	Are clients protected to an acceptable baseline?
Cash and resilience	DSO, customer concentration, vendor concentration, contract renewal profile.	Is the business financially resilient?

METHODOLOGY

Source notes and caveats

This expanded playbook is based on the provided MSP benchmark research and planning note. It consolidates public and partner-summary benchmark material from recognised MSP operating frameworks and market references, including Service Leadership Index references, TruMethods/TruPeer concepts, Channel Futures MSP 501 commentary, MSP industry surveys, pricing calculators, compensation benchmarks and UK/EU regulatory context.

• Several MSP benchmark datasets are paywalled or survey-based, so ranges should be treated as directional.

• Some retention, CAC, LTV:CAC and eNPS ranges use B2B SaaS benchmarks as the closest available proxy.

• UK and Ireland-specific MSP data is thinner than US data, so GBP ranges are rounded and market-adjusted.

• Different sources define top quartile differently: EBITDA, growth, recurring revenue, NPS or composite performance.

• Accounting treatment matters. Gross margin comparisons are weak if one MSP includes service labour in COGS and another does not.

• Security and regulatory requirements change over time, so compliance-specific metrics should be validated against current guidance.

• Benchmark values should support discussion, not replace detailed financial analysis or professional advice.

• This version is intentionally anonymous and vendor-neutral. Vendor names are avoided except where they are part of benchmark source context.

Scorecard note: The benchmark ranges in this playbook represent mature, well-scaled MSP performance and are intentionally aspirational. The accompanying Benchmark Scorecard uses size-adjusted thresholds — calibrated separately for small, mid-market and larger MSPs — rather than a single band applied to all. This is intentional: a 10-person MSP and a 100-person operator work under fundamentally different economics, and applying identical thresholds to both would make the comparison misleading rather than useful. The playbook ranges represent what mature, scaled operators achieve; the scorecard calibrates green, amber and red against what is genuinely strong for each size band, so smaller operators are not penalised for ratios that only larger, more automated businesses can typically reach.